

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.svm import SVC, LinearSVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn import preprocessing

```

```
df = pd.read_csv('Churn_Modelling.csv')
```

```
df.info()
```

```
df.head()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   RowNumber             10000 non-null  int64
 1   CustomerId            10000 non-null  int64
 2   Surname                10000 non-null  object
 3   CreditScore            10000 non-null  int64
 4   Geography              10000 non-null  object
 5   Gender                 10000 non-null  object
 6   Age                   10000 non-null  int64
 7   Tenure                 10000 non-null  int64
 8   Balance                10000 non-null  float64
 9   NumOfProducts          10000 non-null  int64
10   HasCrCard              10000 non-null  int64
11   IsActiveMember         10000 non-null  int64
12   EstimatedSalary        10000 non-null  float64
13   Exited                 10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenu
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	

```
df.drop(columns=['RowNumber', 'CustomerId', 'Surname'], inplace=True)
```

```
df.isna().sum()
```

```
CreditScore      0
Geography        0
Gender           0
Age             0
Tenure           0
Balance          0
NumOfProducts   0
HasCrCard        0
IsActiveMember   0
EstimatedSalary  0
Exited           0
dtype: int64
```

```
df.describe()
```

	CreditScore	Age	Tenure	Balance	NumOfProduct
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	76485.889288	1.5302
std	96.653299	10.487806	2.892174	62397.405202	0.5816
min	350.000000	18.000000	0.000000	0.000000	1.0000
25%	584.000000	32.000000	3.000000	0.000000	1.0000
50%	652.000000	37.000000	5.000000	97198.540000	1.0000
75%	718.000000	44.000000	7.000000	127644.240000	2.0000
max	850.000000	92.000000	10.000000	250898.090000	4.0000

```
X=df.iloc[:, :df.shape[1]-1].values      #Independent Variables
y=df.iloc[:, -1].values                  #Dependent Variable
X.shape, y.shape
```

```
print(X[:8,1], '... will now become: ')
```

```
['France' 'Spain' 'France' 'France' 'Spain' 'Spain' 'France' 'Germany'] ...
```

```
label_X_country_encoder = LabelEncoder()
X[:,1] = label_X_country_encoder.fit_transform(X[:,1])
print(X[:8,1])
```

```
print(X[:6,2], '... will now become: ')
label_X_gender_encoder = LabelEncoder()
X[:,2] = label_X_gender_encoder.fit_transform(X[:,2])
```

```

print(X[:6,2])

[0 2 0 0 2 2 0 1]
['Female' 'Female' 'Female' 'Female' 'Female' 'Male'] ... will now become:
[0 0 0 0 0 1]

transform = ColumnTransformer([("countries", OneHotEncoder(), [1])], remainder="
X = transform.fit_transform(X)
X

array([[1.0, 0.0, 0.0, ..., 1, 1, 101348.88],
       [0.0, 0.0, 1.0, ..., 0, 1, 112542.58],
       [1.0, 0.0, 0.0, ..., 1, 0, 113931.57],
       ...,
       [1.0, 0.0, 0.0, ..., 0, 1, 42085.58],
       [0.0, 1.0, 0.0, ..., 1, 0, 92888.52],
       [1.0, 0.0, 0.0, ..., 1, 0, 38190.78]], dtype=object)

X = X[:,1:]
X.shape

(10000, 11)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

sc=StandardScaler()
X_train[:,np.array([2,4,5,6,7,10])] = sc.fit_transform(X_train[:,np.array([2,4,5
X_test[:,np.array([2,4,5,6,7,10])] = sc.transform(X_test[:,np.array([2,4,5,6,7,1

sc=StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
X_train

array([[ -0.5698444 ,  1.74309049,  0.16958176, ...,  0.64259497,
        -1.03227043,  1.10643166],
       [ 1.75486502, -0.57369368, -2.30455945, ...,  0.64259497,
         0.9687384 , -0.74866447],
       [-0.5698444 , -0.57369368, -1.19119591, ...,  0.64259497,
        -1.03227043,  1.48533467],
       ...,
       [-0.5698444 , -0.57369368,  0.9015152 , ...,  0.64259497,
        -1.03227043,  1.41231994],
       [-0.5698444 ,  1.74309049, -0.62420521, ...,  0.64259497,
         0.9687384 ,  0.84432121],
       [ 1.75486502, -0.57369368, -0.28401079, ...,  0.64259497,
        -1.03227043,  0.32472465]])

# **Initialize & build the model**INPUT = Number columns (Independet ) HIDDEN -

```

```
from tensorflow.keras.models import Sequential
# Initializing the ANN
classifier = Sequential()

from tensorflow.keras.layers import Dense
# The amount of nodes (dimensions) in hidden layer should be the average of input
# This adds the input layer (by specifying input dimension) AND the first hidden
classifier.add(Dense(activation = 'relu', input_dim = 11, units=256, kernel_initializer='uniform'))

# Adding the hidden layer
classifier.add(Dense(activation = 'relu', units=512, kernel_initializer='uniform'))
classifier.add(Dense(activation = 'relu', units=256, kernel_initializer='uniform'))
classifier.add(Dense(activation = 'relu', units=128, kernel_initializer='uniform'))

# Adding the output layer
# Notice that we do not need to specify input dim.
# we have an output of 1 node, which is the the desired dimensions of our output
# We use the sigmoid because we want probability outcomes
classifier.add(Dense(activation = 'sigmoid', units=1, kernel_initializer='uniform'))

# Create optimizer with default learning rate
# sgd_optimizer = tf.keras.optimizers.SGD()
# Compile the model
classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

classifier.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	3072
dense_1 (Dense)	(None, 512)	131584
dense_2 (Dense)	(None, 256)	131328
dense_3 (Dense)	(None, 128)	32896
dense_4 (Dense)	(None, 1)	129

```
=====
Total params: 299009 (1.14 MB)
Trainable params: 299009 (1.14 MB)
Non-trainable params: 0 (0.00 Byte)
```

```
classifier.fit(  
X_train, y_train,  
validation_data=(X_test,y_test),  
epochs=20,  
batch_size=32  
)
```

```
Epoch 1/20  
250/250 [=====] - 4s 9ms/step - loss: 0.4234 - acc  
Epoch 2/20  
250/250 [=====] - 2s 8ms/step - loss: 0.3626 - acc  
Epoch 3/20  
250/250 [=====] - 2s 9ms/step - loss: 0.3500 - acc  
Epoch 4/20  
250/250 [=====] - 3s 11ms/step - loss: 0.3430 - ac  
Epoch 5/20  
250/250 [=====] - 3s 12ms/step - loss: 0.3395 - ac  
Epoch 6/20  
250/250 [=====] - 3s 10ms/step - loss: 0.3337 - ac  
Epoch 7/20  
250/250 [=====] - 2s 8ms/step - loss: 0.3294 - acc  
Epoch 8/20  
250/250 [=====] - 2s 8ms/step - loss: 0.3272 - acc  
Epoch 9/20  
250/250 [=====] - 2s 9ms/step - loss: 0.3242 - acc  
Epoch 10/20  
250/250 [=====] - 4s 14ms/step - loss: 0.3243 - ac  
Epoch 11/20  
250/250 [=====] - 2s 8ms/step - loss: 0.3188 - acc  
Epoch 12/20  
250/250 [=====] - 2s 9ms/step - loss: 0.3138 - acc  
Epoch 13/20  
250/250 [=====] - 3s 11ms/step - loss: 0.3113 - ac  
Epoch 14/20  
250/250 [=====] - 2s 9ms/step - loss: 0.3094 - acc  
Epoch 15/20  
250/250 [=====] - 3s 11ms/step - loss: 0.3030 - ac  
Epoch 16/20  
250/250 [=====] - 2s 8ms/step - loss: 0.2989 - acc  
Epoch 17/20  
250/250 [=====] - 3s 12ms/step - loss: 0.2988 - ac  
Epoch 18/20  
250/250 [=====] - 3s 14ms/step - loss: 0.2940 - ac  
Epoch 19/20  
250/250 [=====] - 5s 19ms/step - loss: 0.2853 - ac  
Epoch 20/20  
250/250 [=====] - 3s 13ms/step - loss: 0.2780 - ac  
<keras.src.callbacks.History at 0x786149a80af0>
```

```
y_pred = classifier.predict(X_test)  
y_pred
```

```
63/63 [=====] - 0s 2ms/step  
array([[0.43192458],  
       [0.21820262].
```

```

.....',
[0.10865229],
...,
[0.22263762],
[0.25179225],
[0.2263274 ]], dtype=float32)

```

```

# To use the confusion Matrix, we need to convert the probabilities that a custo
# So we will use the cutoff value 0.5 to indicate whether they are likely to exi
y_pred = (y_pred > 0.5)
y_pred

```

```

array([[False],
       [False],
       [False],
       ...,
       [False],
       [False],
       [False]])

```

```

##Print the Accuracy score and confusion matrix

```

```

from sklearn.metrics import confusion_matrix, classification_report
cm1 = confusion_matrix(y_test, y_pred)
cm1

```

```

array([[1486, 109],
       [ 185, 220]])

```

```

print(classification_report(y_test, y_pred))

```

```

accuracy_model1 = ((cm1[0][0]+cm1[1][1])*100)/(cm1[0][0]+cm1[1][1]+cm1[0][1]+cm1
print (accuracy_model1, '% of testing data was classified correctly')

```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	1595
1	0.67	0.54	0.60	405
accuracy			0.85	2000
macro avg	0.78	0.74	0.75	2000
weighted avg	0.84	0.85	0.85	2000

```

85.3 % of testing data was classified correctly

```

